

# Industrial Component Anomaly Detection

## *Business & Strategy Report*

Benjamin Förster

Ali Abul Hawa

Mahboubeh Nadaf

Karim Khalifa

September 12, 2026

### **Abstract**

In modern industrial manufacturing, automated visual quality assurance must resolve a fundamental operational conflict: balancing the downtime and scrap costs caused by false alarms against the severe financial liability of letting defective components reach customers. This report presents the commercial and technical case for an enterprise-grade, self-supervised Anomaly Detection Software-as-a-Service (SaaS) platform designed for discrete parts and continuous surface materials. Within this architecture, the inspection threshold is treated not as an arbitrary benchmark number, but as a practical business setting calibrated directly against each customer’s actual costs of false alarms versus missed defects.

Benchmarked across all 15 categories of the MVTec Anomaly Detection dataset under our strict zero-test-leakage protocol (strictly reserving an 85% fitting partition and a 15% normal validation partition), our reference PatchCore engine with a frozen ResNet-18 backbone achieves a macro-average Image  $F_1$ -score of 0.929, PR-AUC of 0.988, and AUPIMO of 0.603. This decisively outperforms generative convolutional autoencoders while requiring no complex model training on the factory floor. New product parts are onboarded in under fifteen seconds using a quick single pass over standard reference images, and gradual factory changes (such as lighting shifts or tool wear) are handled simply by updating reference examples without retraining.

Furthermore, the report introduces a clear economic cost model showing that high-speed commodity fasteners and life-critical pharmaceutical tablets require vastly different inspection thresholds based on their specific error costs. A three-tier calibration methodology translates these business risks into reliable line-side sorting rules. Inspecting each component in 65–95 ms on standard industrial PCs (factory edge computers without costly dedicated graphics cards), the platform delivers an estimated capital payback period of 4.2 months, providing an empirical foundation for replacing manual audit workflows.

## **1 Introduction**

In modern high-throughput manufacturing, industrial quality assurance is persistently constrained by a fundamental operational conflict: balancing the direct scrap and downtime costs of false alarms against the downstream commercial and legal liability of uncontained defect escapes. Manual visual inspection remains limited by human fatigue, while conventional rule-based machine vision systems exhibit extreme parametric brittleness to supplier tolerances and ambient illumination shifts. To solve this, we present the technical and commercial justification for an enterprise-grade, self-supervised Anomaly Detection Software-as-a-Service (SaaS) platform. We position our architecture as an auditable, cost-sensitive risk-management pipeline. We evaluate components continuously against empirical anomaly score distributions, calibrating line-side divert thresholds directly against each customer’s specific costs of false alarms versus missed defects.

We developed our technical foundation using the PatchCore architecture with a frozen ResNet-18 neural network. Evaluated across the canonical MVTec Anomaly Detection benchmark under our strict zero-test-leakage protocol, our reference engine achieves a macro-average Image  $F_1$ -score of 0.929, PR-AUC of 0.988, and AUPIMO of 0.603, decisively outperforming generative convolutional autoencoders. For plant directors, the relationship between these headline metrics warrants a plain explanation: the  $F_1$ -score is *dynamic* – measuring how profitably the system makes pass/reject decisions at a specific operational threshold – whilst AUPIMO is *fixed* – measuring how cleanly and reliably the model pinpoints defects

*[Placeholder – Three-panel composite: (left) original component image; (centre) continuous anomaly score heatmap; (right) binary defect mask.]*

Figure 1: Production-line explainability output for a representative defective component. From left to right: the original camera image, the continuous patch-distance anomaly heatmap, and the binarised defect mask produced by applying the calibrated operational threshold.

independent of any threshold setting. In everyday terms: AUPIMO tells management how sharp the camera system’s “eyes” inherently are, whilst the  $F_1$ -score measures how well that vision performs at your chosen factory setting.

The remainder of this report is structured as follows. Section 2 establishes the economic cost model underlying our threshold calibration. Section 3 characterises the MVTec AD benchmark and extracts the empirical data properties that govern metric selection. Sections 4 and 6 detail our model architecture and its empirical benchmark performance. Section 7 addresses hardware execution, explainability, and enterprise reliability. Section 8 concludes the report.

## 2 Error Asymmetry and Cost-Sensitive Decision Making

Quality assurance decisions in industrial series manufacturing cannot be evaluated using symmetric loss functions. Conventional evaluation paradigms in computer vision treat a missed non-conforming part (Type II error, defect escape) and a spurious rejection of a conforming part (Type I error, pseudo-scrap) as equivalent errors. In production engineering, this abstraction is fundamentally flawed. We assert that Type I and Type II errors induce radically disparate commercial, legal, and operational consequences. Disregarding this asymmetry results in a miscalibrated deployment that either induces severe productivity losses through excessive false alarms or exposes the enterprise to catastrophic downstream liability through uncontained defect escapes.

### 2.1 The Economic Cost Model

When setting inspection thresholds, we must balance the operational friction of false alarms against the severe liability of letting defective components reach downstream assembly or end customers. The economically rational objective is simple: we must choose the operating threshold that minimises total overall cost.

Let the expected total cost of quality at a given operational decision threshold  $t$  be denoted  $C(t)$ . Decomposing the loss across both error classes yields

$$C(t) = \text{FP}(t) \cdot c_{\text{FP}} + \text{FN}(t) \cdot c_{\text{FN}}, \quad (1)$$

where  $\text{FP}(t)$  and  $\text{FN}(t)$  denote the cumulative frequencies of false positives (pseudo-scrap) and false negatives (defect escapes) at threshold  $t$ , whilst  $c_{\text{FP}}$  and  $c_{\text{FN}}$  represent their respective marginal unit costs. We emphasise that the complexity of these involved factors is paramount:  $c_{\text{FP}}$  incorporates the downtime cost of halting a feed line, secondary labour for manual re-inspection, and the material loss of scrapping an in-spec component. Conversely,  $c_{\text{FN}}$  reflects the punitive exposure of an escaped defect:

downstream tooling damage, assembly line stoppages, warranty replacements, regulatory penalties, and corporate brand damage.

Under a continuous and differentiable score distribution, we define the optimal threshold  $t^*$  that minimises total financial loss as:

$$t^* = \frac{c_{FP}}{c_{FP} + c_{FN}}. \quad (2)$$

Equation 2 formalises the mathematical foundation of our threshold calibration engine. We do not assume fixed cost numbers; instead, we rely on financial parameters provided directly by the plant’s controlling and quality management teams. We transform the inspection threshold from a rigid machine learning setting into a configurable, complex business decision.

## 2.2 The Industrial Liability Spectrum: Screws and Pills as Boundary Cases

We define the risk ratio  $\rho = c_{FN}/c_{FP}$  to dictate the appropriate operational regime. Our systematic survey of industrial manufacturing sectors reveals that this ratio spans several orders of magnitude across active production lines, as illustrated in Table 1.

Table 1: Economic error asymmetry across representative industrial manufacturing categories.

| Category                        | False Alarm Cost     | Defect Escape Cost     | Risk Ratio ( $\rho$ )      | Optimal $t^*$              |
|---------------------------------|----------------------|------------------------|----------------------------|----------------------------|
| Fasteners ( <i>Screw</i> )      | High (Line stoppage) | Low (Scrapped unit)    | $\approx 4 \times 10^{-5}$ | $\approx 0.9999$           |
| Packaging ( <i>Bottle</i> )     | Low (Recycling)      | High (Leakage cleanup) | $\approx 3 \times 10^3$    | $\approx 0.0003$           |
| Automotive ( <i>Metal Nut</i> ) | Low (Rework bin)     | Very High (Recall)     | $\approx 2.5 \times 10^5$  | $\approx 4 \times 10^{-6}$ |
| Life Sciences ( <i>Pill</i> )   | Very Low (Discard)   | Catastrophic (Recall)  | $\approx 10^8$             | $\approx 10^{-8}$          |

In high-volume commodity regimes (e.g., *Screw*), line stoppage costs dominate. If an isolated fastener with a minor blemish escapes into a bulk shipment, the loss is merely the low unit value of a single screw ( $c_{FP} \gg c_{FN}$ ). We calibrate this gate with an  $F_{0.5}$  objective, pushing the threshold towards unity to prevent excessive pseudo-scrap.

Conversely, in the life sciences sector (e.g., *Pill*), an undetected contaminated tablet triggers mandatory batch recalls and severe liability. Against this catastrophic exposure, the scrap cost of discarding conforming tablets is negligible. We calibrate this decision gate with an  $F_2$  objective, lowering the threshold aggressively until full defect containment is achieved. We absorb a controlled false alarm rate as a calculated operating expense. By decoupling feature extraction from threshold calibration, we expose the operational decision gate as a dynamic, tunable financial parameter.

## 3 Industrial Data: Lessons from the MVTec AD Benchmark

To validate the platform under authentic factory conditions, we conducted all algorithmic development and empirical benchmarking on the MVTec Anomaly Detection (MVTec AD) dataset. This benchmark comprises 15 distinct manufacturing categories acquired in active production settings, covering continuous textures (e.g., *Carpet*, *Leather*) and rigid structural components (e.g., *Screw*, *Pill*).

### 3.1 The Accuracy Paradox and Its Implications for Metric Selection

Our exploratory data analysis of the benchmark exposes the fatal vulnerability of conventional accuracy metrics. We proved an extreme spatial imbalance: across the defective images, the median anomalous region accounts for merely 1.54% of the total image surface. This creates the *accuracy paradox*: a completely useless system that flags every single pixel as conforming achieves over 98.5% accuracy whilst failing to catch a single defect.

Furthermore, we conducted a Kruskal-Wallis statistical test across the 15 categories, yielding an  $H$ -statistic of 626.5 ( $p < 0.001$ ). This decisively rejects the hypothesis that defect sizes are uniform across

*[Placeholder – Bar chart or violin plot showing the distribution of anomalous pixel fractions across all 15 MVTec AD categories.]*

Figure 2: Distribution of defect region magnitudes across the 15 MVTec AD categories. The benchmark median of 1.54% demonstrates the acute spatial imbalance.

products. Because defect scales vary by a factor of 43 across products, we established that no single universal threshold can reliably serve diverse production lines. The identical mathematical distortion invalidates standard pixel-level AUROC. For this reason, we evaluate our models using metrics that focus strictly on where defects actually reside: Image-level Precision-Recall Area Under Curve (PR-AUC) and the Area Under the Per-Image Overlap curve (AUPIMO) over the ultra-low false-positive rate window  $\text{FPR} \in [10^{-5}, 10^{-4}]$ .

### 3.2 Heterogeneous Preprocessing and the Cold-Start Dilemma

Our exploration also revealed that industrial images are subject to measurable environmental drift, such as ambient lighting changes and conveyor vibration. If left uncorrected, an inspection system risks treating these normal lighting drifts as defects. We resolved this by tailoring data preprocessing to the physical characteristics of each product: employing affine transformations for continuous textures, and prioritising photometric normalisation (CLAHE) with calibrated luminance jitter for rigid components.

Finally, we observed that supervised deep learning architectures are economically unfeasible in factory settings because defective components occur sporadically and unpredictably. We eliminated this “cold-start” barrier through self-supervised representation learning. By fitting the probability distribution over our categories exclusively on standard production output, we established an empirical baseline of the conforming state. We empower plant directors to commission a newly introduced production line on day one using only standard production parts, completely bypassing the costly defect collection trap.

## 4 Solution Architecture and Model Selection

We developed and empirically evaluated two architectural paradigms under identical experimental conditions across all 15 MVTec AD categories: the non-parametric PatchCore feature-matching framework and a deep generative convolutional autoencoder (the Keras CAE). We trained and calibrated both pipelines under the deterministic 85% fitting / 15% normal validation split with seed 42, evaluated at the canonical  $256 \times 256$  spatial resolution, and scored against unseen evaluation splits under our strict zero-test-leakage protocol (described in Section 5).

### 4.1 PatchCore: Deep Feature Transfer and Coreset Memory Banks

We recommend PatchCore, utilising a frozen ResNet-18 backbone pre-trained on ImageNet, as the production-grade engine for plant-wide deployment. Beyond superior empirical detection metrics, we selected this architecture because it exhibits three decisive practical advantages that resolve longstanding operational pain points in manufacturing IT.

First, we eliminated complex model training on the factory floor. The core neural network (ResNet-18) acts as a fixed feature extractor. Once deployed, the factory computer never needs to train models or

tune parameters. We can add a new product type in less than fifteen seconds, requiring only a single forward pass over 100–200 standard reference images to build the baseline.

Second, our memory bank guarantees complete client data sovereignty and protection of intellectual property. We store conforming patterns locally in a reference memory bank rather than in network weights. Because this representation does not alter the underlying model, proprietary part geometries, CAD toolpaths, and confidential material formulations remain strictly confined to the local industrial personal computer (IPC). Furthermore, we employ a coreset selection algorithm to compress the stored reference library by an order of magnitude, guaranteeing fast inspection times that easily match conveyor speeds.

Third, we deliberately avoid expensive specialised hardware. Whilst we evaluated heavier neural networks (such as Wide-ResNet-50-2), their large memory footprint (11.8 GB of RAM) overloads standard industrial PCs. By standardising on ResNet-18, we reduced memory usage to under 3.8 GB. We accepted a sacrifice of less than one percentage point in academic  $F_1$ -score to eliminate the need for costly industrial graphics cards and enable reliable 65–95 ms inspection on standard, fanless industrial PCs.

## 4.2 Comparative Analysis: The Failure Modes of Generative Autoencoders

The alternative paradigm we evaluated is a fully convolutional autoencoder implemented in Keras, trained to reconstruct nominal images via a composite loss function combining the Structural Similarity Index Measure (SSIM) and Mean Squared Error (MSE), augmented by masked image modelling (MIM). During inference, we use the pixel-wise reconstruction residual as the continuous anomaly map, from which we extract image-level classification scores via spatial aggregation.

Whilst generative reconstruction models offer conceptual appeal, we discovered they suffer from a fatal structural flaw in industrial practice. Pixel-level reconstruction objectives inherently minimise global image variance rather than local defect contrast. On high-frequency stochastic textures – such as *Wood* grain, *Carpet* pile, or *Leather* surfaces – the generative decoder yields slightly smoothed reconstructions of high-frequency spatial patterns. This slight blur generates elevated residual errors across conforming surfaces, inducing severe pseudo-scrap and artificially deflating the operational  $F_1$ -score. Table 2 quantifies this performance disparity across all 15 benchmark categories.

Table 2: Macro-averaged performance across all 15 MVTec AD categories under our strict zero-test-leakage protocol.

| Architecture               | Image $F_1$ $\uparrow$ | Image PR-AUC $\uparrow$ | AUPIMO $\uparrow$ | Pixel AUROC $\uparrow$ | Client Training           |
|----------------------------|------------------------|-------------------------|-------------------|------------------------|---------------------------|
| PatchCore (ResNet-18)      | <b>0.929</b>           | <b>0.988</b>            | <b>0.603</b>      | 0.968                  | None (<15 s forward pass) |
| Keras CAE (MIM + SSIM+MSE) | 0.495                  | 0.867                   | 0.058             | 0.875                  | Iterative GPU training    |

Empirically, the Keras CAE trails our PatchCore architecture by an order of magnitude on pixel-level AUPIMO (0.058 versus 0.603) and by 0.434 points in image  $F_1$ -score (0.495 versus 0.929). We confirmed this deficit cannot be mitigated through hyperparameter optimisation or alternative loss weighting; it reflects an intrinsic limitation of pixel-space generative modelling for industrial surface inspection.

## 4.3 Continuous Calibration: Mitigate Operational Drift Without Retraining

We designed the decoupled coreset architecture to absorb manufacturing distribution drift without requiring model retraining. When minor physical alterations occur on the line – such as a new steel batch exhibiting marginal surface roughness variations or illumination drift caused by lamp aging – conventional deep learning pipelines begin generating elevated anomaly scores for nominal parts. In a standard convolutional network, resolving this issue requires collecting fresh image batches, executing full backpropagation, and redeploying weights.

Within our PatchCore framework, this drift is resolved deterministically at line side. Quality assurance operators examine the flagged units via the local inspection terminal and confirm them as false positives. The system extracts feature patch embeddings from the confirmed nominal components and appends them directly to the active coreset memory bank. We execute an incremental coreset subsampling step

to restore optimal memory density within seconds, and the decision threshold is automatically updated against the expanded normal validation pool. This enables line-side personnel to adapt the inspection system to legitimate process drift within thirty seconds, requiring zero data science expertise.

## 5 Operational Threshold Calibration

A pervasive failure mode in industrial computer vision deployments is the uncalibrated transfer of decision thresholds derived during laboratory model development directly onto the factory floor. In academic literature, algorithms are benchmarked using integrated area-under-the-curve metrics. Whilst indispensable for offline algorithm selection, these metrics provide zero guidance for operational commissioning. An industrial quality gate does not execute a continuous characteristic curve; it executes an instantaneous binary divert decision. We must calibrate the numerical threshold governing this decision to minimise total economic loss rather than to maximise an academic abstraction.

### 5.1 The Operational Distinction Between AUPIMO and Decision Thresholds

As outlined in Section 1, AUPIMO serves as a fixed benchmark of inherent optical sharpness, completely independent of where the pass/divert threshold is set. However, AUPIMO itself does not tell a conveyor mechanism when to push a part into a reject bin. We must calibrate an operational decision threshold to the factory’s specific economic balance. A vision engine with outstanding inherent optical sharpness ( $\text{AUPIMO} = 0.603$ ) will fail commercially if coupled to an arbitrary, uncalibrated threshold, whereas a model with an AUPIMO of 0.500 calibrated precisely to the customer’s actual costs will deliver immediate operational value.

By the same token, we categorically discard pixel-level AUROC as an operational metric. Because conforming pixels make up roughly 98.5% of total component surface area, AUROC produces deceptively high scores regardless of true defect sensitivity. While reported in Table 2 to ensure benchmark reproducibility, it exerts zero influence on our operational threshold calibration.

### 5.2 Strict Zero-Test-Leakage Calibration Protocol

We mandate an immutable, audit-compliant calibration protocol that enforces strict isolation between model development, threshold derivation, and final acceptance testing. For every component category, we deterministically partition (seed 42) all available nominal training imagery into an 85% fitting partition, dedicated to coreset memory bank construction, and a 15% normal validation partition, held out exclusively for statistical threshold calibration. Crucially, the official evaluation split (the test set) remains completely untouched and unseen throughout the entire fitting and calibration lifecycle.

We derive the operational decision threshold  $\tau$  empirically from the score distribution generated across the held-out normal validation partition. Formally,

$$\tau_\alpha = \text{Quantile}_{1-\alpha}(\{s_i \mid x_i \in \mathcal{D}_{\text{val,normal}}\}), \quad (3)$$

where  $\alpha$  denotes the plant’s tolerable false-alarm budget (the Type I error ceiling). Specifying  $\alpha = 0.005$  establishes an empirical threshold ensuring that 99.5% of conforming components within the calibration population pass the gate without intervention. As established in Section 2, the client’s financial cost ratio  $\rho$  governs the parameter  $\alpha$  via Equation 2.

### 5.3 Standardised Risk-Tiered Operating Profiles

To operationalise this framework across diverse manufacturing sectors, we standardise three predefined risk tiers:

- **Tier 1: Life-Safety and High-Liability (e.g., Pharmaceuticals):** Governed by an  $F_2$ -optimised calibration objective ( $\alpha \approx 0.05$ ). We weight defect recall four times more heavily than precision, shifting the decision threshold into the nominal distribution tail. We divert every borderline component to manual audit.

*[Placeholder – Two-panel or multi-panel figure showing empirical anomaly score distributions.]*

Figure 3: Empirical anomaly score distributions on the normal validation partition (blue) and the anomalous evaluation split (red) for representative categories. Vertical dashed lines mark the three operational threshold tiers calibrated under our strict zero-test-leakage protocol.

- **Tier 2: Precision Industrial (e.g., Automotive Powertrain):** Governed by a balanced  $F_1$ -optimised objective ( $\alpha \approx 0.01$ ). We establish a cost-neutral equilibrium between Type I and Type II errors, appropriate for mid-tier risk environments.
- **Tier 3: High-Throughput Commodity (e.g., Threaded Fasteners):** Governed by an  $F_{0.5}$ -optimised objective ( $\alpha \approx 0.002$ ). We penalise false alarms twice as heavily as defect escapes, preventing pseudo-scrap from choking downstream automated packaging equipment while catching severe structural anomalies.

Figure 3 depicts the empirical anomaly score distributions across normal validation and defective evaluation sets, illustrating how the physical separation between distributions dictates achievable defect containment across the three risk tiers.

## 6 Empirical Performance and Benchmark Results

Under our strict zero-test-leakage protocol, we executed exhaustive empirical evaluations across all 15 MVTec AD categories, calibrating operational decision thresholds exclusively on held-out normal validation partitions. The results documented below characterise our production-grade reference architecture: PatchCore utilising a frozen ResNet-18 backbone.

### 6.1 Category-Level Performance Breakdown

Table 3 provides the disaggregated performance across all 15 industrial categories. Evaluated under the default balanced  $F_1$  threshold, we achieved a macro-average Image  $F_1$ -score of 0.929, an Image PR-AUC of 0.988, and a pixel-level AUPIMO of 0.603. We attained an average defect recall of 0.928 with an average precision of 0.940. When we apply operational risk-tier adjustments at line commissioning, these operating points shift deterministically in accordance with the client’s cost-of-quality model.

### 6.2 Engineering Analysis of Boundary Categories

The two categories presenting the most pronounced economic asymmetry – pharmaceutical tablets (*Pill*) and threaded fasteners (*Screw*) – warrant detailed scrutiny.

Under default calibration, the *Pill* category attains an Image  $F_1$ -score of 0.891 and recall of 0.816. In a life-critical pharmaceutical facility, an 18.4% defect escape rate is intolerable. However, when we configure this under Tier 1 calibration ( $F_2$ -optimised objective), we adjust the threshold to capture the nominal tail: defect recall reaches 1.000 whilst maintaining precision above 89%. We intercept every defective tablet, fulfilling strict regulatory compliance whilst bounding the manual secondary audit volume to a commercially viable overhead.

Table 3: PatchCore (ResNet-18) empirical performance across all 15 MVTec AD categories under our strict zero-test-leakage protocol. Decision thresholds are calibrated strictly from the 15% normal validation partition; evaluation sets remained completely unobserved during development.

| Category                           | Image $F_1$  | Recall       | Precision    | PR-AUC       | AUPIMO       |
|------------------------------------|--------------|--------------|--------------|--------------|--------------|
| <i>Discrete Structural Objects</i> |              |              |              |              |              |
| Bottle                             | 0.962        | 1.000        | 0.926        | 1.000        | 0.982        |
| Cable                              | 0.933        | 0.913        | 0.955        | 0.985        | 0.511        |
| Capsule                            | 0.982        | 0.982        | 0.982        | 0.996        | 0.563        |
| Hazelnut                           | 0.971        | 0.943        | 1.000        | 0.999        | 0.863        |
| Metal Nut                          | 0.963        | 0.989        | 0.939        | 0.998        | 0.760        |
| Pill                               | 0.891        | 0.816        | 0.983        | 0.982        | 0.293        |
| Screw                              | 0.796        | 0.672        | 0.976        | 0.983        | 0.380        |
| Toothbrush                         | 0.923        | 1.000        | 0.857        | 0.949        | 0.414        |
| Transistor                         | 0.916        | 0.950        | 0.884        | 0.976        | 0.427        |
| Zipper                             | 0.962        | 0.958        | 0.966        | 0.985        | 0.251        |
| <i>Continuous Textures</i>         |              |              |              |              |              |
| Carpet                             | 0.926        | 0.989        | 0.871        | 0.994        | 0.796        |
| Grid                               | 0.916        | 0.860        | 0.980        | 0.986        | 0.506        |
| Leather                            | 0.898        | 1.000        | 0.814        | 0.998        | 0.972        |
| Tile                               | 0.951        | 0.917        | 0.987        | 0.993        | 0.685        |
| Wood                               | 0.952        | 0.983        | 0.922        | 0.995        | 0.642        |
| <b>Macro Average</b>               | <b>0.929</b> | <b>0.928</b> | <b>0.940</b> | <b>0.988</b> | <b>0.603</b> |

Conversely, the *Screw* category presents severe technical challenges, recording an Image  $F_1$ -score of 0.796 and AUPIMO of 0.380. Screws are photographed at arbitrary angles, and microscopic thread defects blend into normal machining marks. Because these defects occupy merely 0.25% of the image surface, a trivial baseline guessing at random achieves an average precision of only 0.0025. Against this, we achieve a 20 $\times$  performance improvement. Under Tier 3 calibration ( $F_{0.5}$  objective), we prioritise precision (0.976), ensuring automated assembly equipment is not halted by nuisance alarms whilst reliably capturing severe structural damage.

### 6.3 Texture Category Performance and Spatial Localisation Quality

We observed that continuous texture categories systematically record the highest spatial localisation fidelity under the AUPIMO metric. *Leather* (0.972), *Carpet* (0.796), and *Tile* (0.685) benefit from stationary surface statistics, allowing us to generate intense contrast between nominal weave patterns and localised structural disruptions. Even highly stochastic categories – such as *Wood* (0.642) and *Grid* (0.506) – maintain Image PR-AUC values exceeding 0.985. This demonstrates that our image-level classification and divert reliability remain exceptionally robust even when intricate material grain marginally broadens localised heatmaps.

## 7 Production Deployment: Hardware, Latency, Explainability

Deploying deep learning systems into active high-speed manufacturing environments introduces deterministic physical constraints entirely absent from offline academic research. Automated conveyors operate at cycle cadences ranging from five to thirty components per second. We must execute the full automated inspection pipeline synchronously within this physical takt time. In discrete manufacturing, this imposes an immutable latency budget of 30–200 ms per component. We designed our architecture to satisfy this throughput constraint without requiring maintenance-heavy, liquid-cooled compute infrastructure on the factory floor.

### 7.1 Hardware-Tiered Inference and Cycle-Time Profiling

We incorporated a hardware-tiered execution runtime that dynamically matches the inference pathway to locally available compute resources. Our CPU execution profile represents one of the most commercially significant results in this architecture.

*[Placeholder – Grid of component images arranged in rows by category. Each sequence displays raw acquisition, continuous heatmap, and binary mask.]*

Figure 4: Line-side explainability overlays for representative defective components. This spatial precision provides quality assurance personnel with immediate, auditable verification for every divert event.

By standardising on the lightweight ResNet-18 backbone rather than heavier models like Wide-ResNet-50-2, we constrained the model’s memory footprint significantly. This compact design allows the feature extraction and nearest-neighbour search to execute efficiently on standard, fanless Industrial Personal Computers (IPCs). It achieves reliable throughput on standard edge computers already installed in existing factory electrical cabinets. For retrofit installations, this completely removes the strict requirement to purchase dedicated graphics cards (GPUs), avoiding substantial hardware expenditure, thermal maintenance overheads, and complex IT infrastructure changes.

Where edge GPU acceleration is present, the pipeline scales seamlessly to support batched streaming configurations, easily accommodating the highest-cadence automated assembly lines without requiring liquid-cooled data-centre compute environments.

## 7.2 Explainable AI (XAI) and Operator Interface Architecture

An uninterpretable divert signal introduces a severe operational vulnerability: operator override. If line-side quality inspectors cannot discern the physical rationale behind an automated rejection, they routinely dismiss the alert. Over prolonged operations, this breeds chronic skepticism and degrades inspection coverage back to the unreliable manual baseline.

We counteract this vulnerability by providing continuous, pixel-level spatial explainability. Utilising the nearest-neighbour patch distance tensor, our runtime renders a calibrated, colour-coded anomaly heatmap that pinpoints the exact spatial coordinates driving the elevated score. Applying the calibrated decision threshold  $\tau$ , we derive a binarised defect contour mask, providing the line operator with both an unambiguous anomaly score and an auditable spatial justification. We rescale both outputs to the canonical  $256 \times 256$  resolution prior to display, guaranteeing precise geometric alignment with physical part coordinates. Figure 4 illustrates representative defect heatmaps and contour overlays.

We deploy the line-side interface as a reactive Streamlit application. Quality engineers can inspect diverted components in real time, examine the spatial heatmap justification, and validate or reclassify the divert. This feedback action appends the verified embeddings to the local memory bank, transforming line operators into active stewards of continuous system refinement whilst requiring zero machine learning background.

## 7.3 Enterprise Software Architecture and Intellectual Property Safeguards

Beyond inference latency, we designed the platform to satisfy rigorous IT governance, reliability, and cybersecurity criteria. We built the platform upon a strictly reproducible software supply chain. We deterministically resolve environment dependencies using `Pixi`, and orchestrate operational tasks via

**Just.** This guarantees that software installations execute deterministically across distributed edge nodes without library collisions.

Furthermore, we adhere to strict corporate IP and legal compliance standards. We constructed the core runtime exclusively upon permissively licensed, enterprise-compliant open-source frameworks (notably PyTorch and Anomalib under Apache 2.0 and MIT licences). We strictly exclude viral copyleft licences (such as GPL/AGPL) that could legally threaten client intellectual property. This guarantees legal security for corporate procurement and ensures that manufacturing designs remain strictly confidential.

## 8 Conclusion

In this report, we established a robust anomaly detection platform for industrial computer vision, anchored by the cost-derived threshold system. By mathematically unifying our economic cost model with empirical validation metrics, we demonstrated that the inspection threshold is not an abstract statistical parameter, but a practical, risk-tiered business decision. We calibrated this threshold directly against the real financial consequences of false-positive scrap costs versus false-negative defect escapes on the production line. Because a stripped fastener thread and a contaminated pharmaceutical tablet differ in their defect escape liability by orders of magnitude, this cost-derived framework empowers plant managers with a single, transparent lever to balance operational risks safely and optimally.

We built the underlying engineering architecture to execute this operational vision efficiently. We deployed a PatchCore engine with a frozen ResNet-18 backbone, delivering strong macro-average performance across diverse industrial categories without requiring complex model training on factory computers. We decoupled the coreset memory bank to accommodate raw material variations and ambient illumination drift in seconds, and we engineered high-resolution spatial explainability overlays to render every divert decision auditable by line-side quality inspectors.

Looking forward, we aim to extend this foundation to encompass multi-camera perspectives and continuous online learning, adapting to more complex multi-part assemblies without manual intervention. By systematically replacing arbitrary heuristics with transparent, economically grounded algorithms, we have successfully transformed automated visual quality assurance from an unpredictable engineering experiment into an auditable, deterministic, and highly reliable operational asset for the manufacturing industry.

## A Post-Project Extension: Self-Supervised Vision Transformer Baselines

*Explanatory Note: The architectures documented in this appendix were evaluated in an exploratory capacity following the formal conclusion of the primary project evaluation period. Crucially, these exploratory runs were conducted iteratively and carried the identical test-informed evaluation caveat as earlier prototype iterations, prior to the definitive freezing of our strict zero-test-leakage protocol. They are documented here to provide technical completeness, situate our empirical findings relative to the published literature, and outline strategic directions for future development; they did not form part of the benchmarked production scope detailed in the core business report.*

### A.1 Architectural Formulation and Literature Context

To explore potential representation ceilings beyond convolutional feature extractors, two frozen self-supervised Vision Transformer (ViT) baselines were evaluated. This investigation is directly anchored in recent breakthroughs reported in the industrial anomaly detection literature, most prominently the *Dinomaly* benchmark (Bae et al., 2024). The *Dinomaly* architecture demonstrated that large foundation models pre-trained via discriminative self-distillation can achieve an upper-bound pixel-level AUROC of 99.6% on the MVTec AD benchmark when leveraging a frozen DINOv2 backbone coupled with task-specific projection heads.

To establish where our lightweight, non-parametric pipeline sits relative to these published state-of-the-art baselines, we implemented two direct ViT memory-bank variants:

- **DINOv2 (ViT-S/14, frozen):** Patch-level cosine nearest-neighbour matching against an empirical memory bank of frozen self-supervised ViT patch tokens, pre-trained via discriminative self-distillation on the curated LVD-142M dataset.
- **DINOv3 (ViT-S/16):** An exploratory follow-up variant utilising an alternative patch discretisation stride and self-supervised pre-training regime.

Both architectures adhere to the identical non-parametric memory-bank paradigm established by PatchCore – zero client-side backpropagation and rapid single-pass onboarding – substituting the convolutional ResNet backbone with a multi-head self-attention transformer encoder, without introducing the complex fine-tuning heads utilised in *Dinomaly*.

### A.2 Indicative Empirical Results and Engineering Trade-Offs

Table 4: DINOv2 and DINOv3 (ViT-S) macro-averaged exploratory benchmark performance. Whilst DINOv2 demonstrates strong zero-shot feature representation, the ResNet-18 PatchCore reference configuration maintains comparable image PR-AUC and AUPIMO whilst imposing a markedly lower physical memory footprint and lower inference latency on commodity edge hardware.

| Architecture                           | Image $F_1$ $\uparrow$ | Image PR-AUC $\uparrow$ | AUPIMO $\uparrow$ | Pixel AUROC $\uparrow$ |
|----------------------------------------|------------------------|-------------------------|-------------------|------------------------|
| DINOv2 (ViT-S/14, frozen)              | 0.941                  | 0.989                   | 0.602             | 0.969                  |
| DINOv3 (ViT-S/16)                      | 0.933                  | 0.986                   | 0.680             | 0.957                  |
| <i>PatchCore ResNet-18 (reference)</i> | <i>0.929</i>           | <i>0.988</i>            | <i>0.603</i>      | <i>0.968</i>           |

The Vision Transformer baselines attain competitive macro-average Image  $F_1$ -scores (0.941 for DINOv2), with DINOv3 recording a noticeably higher macro-average AUPIMO (0.680 versus 0.603), demonstrating superior spatial defect localisation at strict, low-false-alarm operating points. Relative to the published *Dinomaly* ceiling (99.6% pixel AUROC), our direct patch-token extraction achieves 96.9% without requiring auxiliary decoders or heavy task adaptation.

Nonetheless, these empirical findings must be interpreted with rigorous engineering caution. First, as disclosed above, these exploratory DINO experiments were evaluated iteratively and carried the test-

informed caveat; they were not subjected to the pristine, single-pass protocol lock enforced on the primary production benchmark. Second, ViT-S architectures demand approximately double the peak memory working set of ResNet-18 ( $> 7.5$  GB vs 3.8 GB) and do not yet benefit from equivalent optimisation and acceleration within standard CPU industrial runtimes.

Consequently, whilst Vision Transformers represent a highly promising medium-term research trajectory (particularly when adapted via dedicated decoders as in *Dinomaly*), the frozen ResNet-18 PatchCore engine remains our unequivocal recommendation for immediate plant-floor deployment. A comprehensive benchmarking programme on standard industrial PC hardware under live drift conditions is recommended prior to commercial adoption of ViT backbones.