

Industrial Component Anomaly Detection

Project Presentation

Unsupervised defect classification and pixel-level anomaly localisation from normal-only manufacturing parts.

Unsupervised AD

Deep Transfer Learning

Zero-Leakage Protocol

TEAM

B. Förster · A. Abul Hawa · M. Nadaf · K. Khalifa

MVTEC AD 2D Dataset

Selected Examples · Objects & Textures

CATEGORY


CONFORMING

DEFECTIVE


Bottle

Broken rim

GOOD



DEFECT



Capsule

Scratch

GOOD




DEFECT




Leather

Surface cut

GOOD




DEFECT




Screw

Thread defect

GOOD



DEFECT



Data Exploration

The MVTec AD Dataset – Class Imbalance

Dataset at a glance

- 15 manufacturing categories: 10 objects + 5 textures
- > 5000 images, $\approx$ 1200 anomalous images
- Normal-only training; test set has pixel-level anomaly ground truth


Accuracy paradox

- Median anomalous-pixel area: **1.54%**
- Predicting *all-normal* yields >98.5% pixel accuracy while catching zero defects.


Kruskal-Wallis: $H(14) \approx 626.5$, $p < 0.001$

- $43\times$ spread in defect size across categories.

TEST-SET CLASS IMBALANCE (images)



PIXEL-AREA IMBALANCE (defect region %)



Defect Taxonomy & Spatial Distribution

Two fundamentally different defect regimes:

Structural Objects (e.g. Screw, Cable)

- Defects cluster at specific locations
- Strong positional bias (shortcut risk)
- Foreground/background separation critical

Continuous Textures (e.g. Grid, Leather)


- Defects broadly distributed across surface
- Rich augmentation needed (rotation, scale)
- No foreground mask required

Covariate shift: Screw, Leather, Grid show measurable brightness/contrast drift between train and test normal images – motivates photometric augmentation.

OBJECTS – LOCALISED DEFECTS (contour = defect boundary)



Screw thread




Cable missing




Capsule scratch

TEXTURES – DISTRIBUTED DEFECTS (defect across full surface)



Grid glue strip



Leather cut



GT mask example


Metric Selection: AUPIMO & F1-Score

The problem with Pixel-AUROC

- AUROC relies on False Positive Rate
 - With $\approx 98.5\%$ normal pixels, TN dominates $\Rightarrow$ FPR barely moves.
 - SOTA scores all exceed 99%, making differences indistinguishable.

Our Evaluation Paradigm

- **AUPIMO:** Pixel-level threshold over the ultra-low FPR window $[10^{-5}, 10^{-4}]$ on normal images.
- **Dual Thresholds:** We calibrate distinct *image-level* and *pixel-level* anomaly thresholds.
- **F1-Score Focus:** Additionally, we always evaluate the F1-Score.




Binarization threshold curve & AUPIMO range

Models

Model Selection & Architecture Overview

1. Convolutional-Autoencoders


Paradigm: Generative Reconstruction · Built from scratch



- **Complete control:** Full freedom over latent bottleneck dimensions, skip connections, and loss constraints.
- **Intuition:** Trained on normal images; unable to reconstruct defects.
- **Trade-offs:** Prone to blurriness or identity shortcuts; sensitive to reconstruction thresholds.

2. PatchCore + ResNet-18 (Anomalib)



Paradigm: Deep Feature Memory Bank · High accessibility



- **High accessibility:** Zero model training required; out-of-the-box integration via Intel's `anomalib`.
- **ResNet-18 backbone:** Compact ImageNet representations; mid-level layers capture both texture and shape.
- **Coreset sampling:** Compresses the memory bank by 90%+ without losing local anomaly detection accuracy.

PatchCore: Transfer Learning & Architecture

Transfer Learning Paradigm



Wood, zipper, transistor, toothbrush. White = GT; green = predicted mask (pixel $q_{0.99}$).

Key Features

- **Cold-start eliminated:** Deployment without defect collection.
- **IP Protection:** Memory bank contains embeddings, not raw model weights.

Why it excels

- Mid-level layers capture local textures perfectly.
- No catastrophic forgetting (frozen backbone).
- ResNet-18 guarantees predictable GPU memory bounds.

Keras CAE – Autoencoder Learning Process

The Concept

Training Phase:

- Trained *excl.* on defect-free images.
- Learns the structural grammar of "normal" surfaces.


Inference Phase:

- On defective images, anomalies are reconstructed as "normal" parts.
- Reconstruction fails at the defect sites $\Rightarrow$ high reconstruction error.
- High reconstruction error $\Rightarrow$ Anomaly Signal.

Key Design Decisions

- **Masked Image Modelling (MIM):**
25% of patches randomly masked before encoding. Forces the model to understand global context, prevents "identity mapping".
- **SSIM + MSE Loss:**
 $\mathcal{L} = \alpha (1 - \text{SSIM}) + (1 - \alpha) \text{MSE}$.
MSE captures global variance, while SSIM enforces high structural and perceptual fidelity.
- **Architectural Choices:**
ELU no dying neurons + zero-centered mean.
BatchNorm stabilizes covariate shifts.
AdamW enables decoupled weight decay.

Keras CAE – End-to-End Pipeline



Automated Pipeline & Optuna Integration

The pipeline is dynamically orchestrated. Hyperparameters (Preprocessing/Model) (e.g., augmentation strategy, latent dimensions) optimized using Optuna targeting **Pixel AUPIMO**.

Evaluation Protocol & Benchmarks

Zero-Leakage Evaluation Protocol

The Core Principle

Every evaluation result in this presentation is produced under a deterministic, audit-compliant split that keeps the test set completely unseen until final scoring $\Rightarrow$ **zero-test-leakage**.


Peak Performance (PatchCore vs CAE)

Category	Metric	PatchCore	Keras CAE
Bottle (Object)	Image F1	0.962	0.992
	AUPIMO	0.982	0.507
Macro Average	Image F1	0.929	0.495
	AUPIMO	0.603	0.058

Canonical Evaluation Settings

- **Split Size:** 85% fitting / 15% validation
- **Test Set:** held out - normal + anomalies
- **Resolution:** 256×256 pixels.
 - Bilinear interpolation for continuous heatmaps.
 - Nearest-neighbor + binarization for GT masks.
- **Threshold:** 95th percentile ($\alpha = 0.05$ false-alarm budget).
- **AUPIMO:** Standard bounds
 $\text{FPR} \in [10^{-5}, 10^{-4}]$, 50k thresholds.

Qualitative Results - Localisation Comparison



Key Observations:

- **PatchCore:** Heatmaps sharply bound the defect, leading to precise predicted masks.
- **Keras CAE:** Reconstruction errors diffuse. While broadly highlighting anomalous regions, they fail at fine-grained pixel localization.

Business Interpretation & Deployment Strategy

Economic Error Asymmetry

A false alarm and a missed defect have different consequences.

False alarm

A conforming part is flagged. Operators may need to inspect it again, slowing the line or creating scrap.

Throughput-sensitive lines may put greater weight on this cost.

Missed defect

A defective part passes inspection.

Downstream rework, warranty claims, recalls or safety risks may follow.

High-consequence applications may put greater weight on this cost.

$$C(\tau) = c_{FP} FP(\tau) + c_{FN} FN(\tau)$$

A decision framework: plant-specific costs and error counts must be measured before choosing a gate.

Threshold Calibration & Operating Priorities

Normal-validation calibration

$$\tau_{\alpha} = \text{Quantile}_{1-\alpha}(\text{AnomalyScores}_{\text{normal, val}})$$


α is a proposed false-alarm budget on held-out normal examples, not a guaranteed production rate.

Proposed priorities

Recall priority ($\alpha \approx 0.10$): lower gate, more review.

Balanced starting point ($\alpha = 0.05$): reported benchmark setting.

Throughput priority ($\alpha \approx 0.01$): higher gate, fewer flags on normal validation data.




Schematic only. Defect recall and real-line false alarms must be checked with plant data for each proposed setting.

Operator Evidence for Review

A flagged image needs visible evidence and a human decision path.



Input image



Anomaly map



Predicted contour

PatchCore highlights suspected defect regions; the heatmap does not establish the cause.

Borderline flags should go to quality engineers for review during a controlled pilot. Review time and operator agreement still need measurement.

Conclusion & Next Steps

What We Demonstrated

- EDA revealed the limitations of conventional evaluation metrics.
- Developed and compared convolutional autoencoders and PatchCore.
- Established a consistent evaluation and threshold-calibration framework.

What Comes Next

- Validate on representative production data.
- Calibrate thresholds against plant-specific risks.
- Run a controlled industrial pilot before deployment.

From understanding defects to building a model, and from model performance to industrial decisions.

Thank you

Questions & Discussion

B. Förster · A. Abul Hawa

M. Nadaf · K. Khalifa